
SAT solving for Python and Rust programmers

A tutorial for people who know boolean logic
and have never used a solver

Carlo Perassi

For [cdclkit](#) 0.1.2 and [dratify](#) 0.1.2

Who this is for. You can read $(a \text{ and not } b) \text{ or } c$ and say whether it is true for given inputs. You have never used a SAT solver, and you are not sure why you would. This assumes nothing else.

What you will be able to do. Turn a real problem — scheduling, colouring, configuration, puzzle — into something a solver answers in milliseconds, and know when the answer can be trusted.

Every code block here was executed, and every output is what it actually printed.

Contents

Part I — The theory you need, and no more	3
1 What a SAT problem is	3
2 Why you would install this	3
3 Conjunctive normal form	4
4 DIMACS, the file format	5
5 The asymmetry that matters	5
6 Proofs of unsatisfiability	5
7 What a solver actually does	6
Part II — One problem, end to end	7
8 The timetable	7
9 The pipeline	7
10 Encoding it	7
11 What actually got produced	9
12 Picking it up from Rust	10
Part III — Python	12
13 Installing	12
14 Your first formula	12
15 Unsatisfiable, and proving it	12
16 Constraints, without writing clauses	13
17 The modelling layer	14
18 Why is it unsatisfiable?	15
19 The command line	15

Part IV — Rust	17
20 What the Rust crate is for	17
21 Literals	17
22 Checking a proof	17
23 What the result tells you	18
24 Two implementations on purpose	19
Where to go next	19

Part I

The theory you need, and no more

1 What a SAT problem is

A *SAT problem* asks one question:

Given a boolean formula, is there any assignment of true/false to its variables that makes the whole thing true?

That is it. If yes, the formula is **satisfiable** (SAT) and the solver hands you an assignment. If no, it is **unsatisfiable** (UNSAT).

This sounds like a toy. It is not, for one reason: an enormous number of practical problems are *disguised* SAT problems. Scheduling, routing, hardware verification, package dependency resolution, timetabling, puzzle solving — all of them reduce to “find an assignment satisfying these constraints”. SAT solvers have become good enough that translating your problem into boolean logic and handing it over often beats writing a bespoke algorithm.

2 Why you would install this

Two different answers, because the two packages are for different jobs. Read the one that applies to you.

If you write Python: **cdclkit**

You reach for this when you have a pile of constraints and need any assignment satisfying all of them. Concretely:

- **Rostering and scheduling.** Six nurses, three shifts, nobody works two nights in a row, everyone gets a weekend off a month.
- **Configuration.** Which package versions can coexist. Which hardware options are mutually compatible.
- **Assignment.** Exams to rooms and time slots with no clashes. Frequencies to transmitters that must not interfere.
- **Puzzles and games.** Sudoku, nonograms, Einstein’s zebra — genuinely the same shape of problem as the three above.

The alternative is writing a backtracking search yourself. That is slow to write, and the version you write in an afternoon will be far slower than a solver that has had thirty years of engineering poured into the same loop. Encoding to SAT hands that part to someone else.

Two things you get that a hand-written search does not:

- **When there is no solution, you learn why.** A minimal unsatisfiable subset (Section 18) names the handful of constraints that actually conflict, rather than reporting “infeasible” and leaving you to bisect a thousand of them by hand.
- **When it says “no”, you can check it.** That is the whole of Section 5, and it is unusual — most tools ask you to take “no solution exists” on faith.

When not to. If your problem is numeric optimisation — maximise profit subject to linear constraints over real numbers — you want a linear-programming or MIP solver, not this. If you need raw throughput on industrial instances with millions of clauses, install `python-sat`, which ships kissat and CaDiCaL as binary wheels and will be far faster than anything here.

So what is actually different? Less than the previous paragraph might suggest, and it is worth being precise. `python-sat` installs just as easily — wheels on every platform, one small pure-Python dependency — so “no dependencies” buys you very little on its own. Its Python layer is perfectly readable too.

One difference is real and checkable: `python-sat` can *emit* DRUP proofs (`with_proof=True`) and ships **nothing to check them**. Its solvers are third-party C and C++ compiled into an extension module, so an UNSAT answer is something you take on trust unless you fetch a separate C tool. Here, checking is on by default and the checker is a package you can read.

Whether that matters depends entirely on what you are doing. If you are searching for a solution, it does not. If you are about to act on “no solution exists”, it is the only thing that does.

If you write Rust: the `dratify` crate

Narrower, and sharper. You install this when you need to **verify a proof of unsatisfiability inside a Rust process**.

That sounds niche until you hit it. Something — a solver in any language, a model checker, a build tool — reported that no solution exists, and you are about to act on that. Ship the hardware. Declare the configuration impossible. Sign off the safety argument. The report is a bare assertion from a large optimised program, and you would like evidence.

Until now your options in Rust were to shell out to `drat-trim` — C you have to compile and invoke as a subprocess — or to trust the solver. This crate is a dependency-free alternative that does RUP and RAT properly and runs in-process.

When not to. If you want to *solve* from Rust, this is not that; use an established solver crate. This checks proofs. The distinction is the point of Section 5.

3 Conjunctive normal form

Solvers do not accept arbitrary formulas. They accept **CNF**: an AND of ORs.

- A **literal** is a variable or its negation: $a, \neg a$.
- A **clause** is an OR of literals: $(a \vee \neg b \vee c)$.
- A **formula** in CNF is an AND of clauses.

$$(a \vee b) \wedge (\neg a \vee b) \wedge (a \vee \neg b)$$

To satisfy the formula you must satisfy *every* clause, and to satisfy a clause you need *at least one* of its literals true. That is the whole semantics.

Is CNF a limitation?

No. Every boolean formula can be converted to CNF, and there is a standard trick (Tseitin encoding) that does it without the exponential blow-up the naive conversion would cause. It works by naming subexpressions: introduce a fresh variable t standing for $(x \wedge y)$, add clauses forcing $t \leftrightarrow (x \wedge y)$, and use t from then on.

You rarely do this by hand. The library does it for you — Section 16 shows the interface.

4 DIMACS, the file format

Everyone in this field exchanges CNF in one plain-text format, and you will see it constantly.

```
p cnf 2 3
 1 2 0
-1 2 0
 1 -2 0
```

- `p cnf 2 3` — header: 2 variables, 3 clauses.
- Variables are numbered from **1**. A negative number means negation.
- Each clause is a list of literals terminated by `0`.
- Lines starting with `c` are comments.

So the file above is exactly $(a \vee b) \wedge (\neg a \vee b) \wedge (a \vee \neg b)$, with $a = 1$ and $b = 2$.

5 The asymmetry that matters

Here is the thing that makes this field interesting, and it is worth pausing on.

A “yes” answer checks itself. If the solver says satisfiable and hands you an assignment, you verify it in linear time: walk the clauses, check each has a true literal. You do not have to trust the solver at all. You do not even have to understand it.

A “no” answer does not. If the solver says unsatisfiable, what has it given you? Nothing you can check. It is asserting that *none* of the 2^n assignments works. Re-deriving that yourself means redoing the search.

This is not paranoia. Modern SAT solvers are tens of thousands of lines of aggressively optimised C, and bugs in them are found regularly — including in solvers that have won competitions. If your hardware verification tool reports “no counterexample exists”, you would like that to mean something.

6 Proofs of unsatisfiability

The field’s answer is that the solver must **show its work**.

While searching, a solver learns new clauses that are consequences of the original formula. If it can eventually derive the **empty clause** — a clause with no literals, which cannot be satisfied by anything — then the formula is unsatisfiable. Logging that sequence of derived clauses is a **DRAT proof**.

A separate program, a *checker*, replays the log and verifies each step follows from what came before. The checker is small enough to audit, and shares no code with the solver.

RUP, informally

The main proof rule is **RUP** (reverse unit propagation). To justify a new clause C :

1. Assume C is false, which forces every literal in it to be false.
2. Propagate: any clause with all-but-one literal false forces its remaining literal.
3. If that produces a contradiction, C must have been true. Accept it.

There is a stronger rule, **RAT**, which permits clauses that are not strictly implied but preserve satisfiability. It is what the “A” in DRAT stands for. You do not need it to use a solver; you need it if you write a checker.

7 What a solver actually does

You do not need this to use one, but a paragraph helps.

Modern solvers use **CDCL** — conflict-driven clause learning. Roughly: pick an unassigned variable and guess a value; propagate the forced consequences; if you hit a contradiction, analyse *why* and derive a new clause that forbids the mistake; backtrack and continue. The learned clauses accumulate, and each one prunes a chunk of the search space. That learning step is what separates CDCL from brute force, and it is why solvers handle problems with millions of variables.

Part II


One problem, end to end

Before the details, the whole journey once. Same problem, both languages, and every file that passes between them printed in full — they are small enough to read, which is the point of choosing a small problem.

8 The timetable


Three exams: Algebra, Biology, Chemistry. Some students are enrolled in both Algebra and Biology, some in both Biology and Chemistry, some in both Algebra and Chemistry. Two exams cannot share a slot if a student sits both.

How many slots do you need?



An edge means “shares a student, so cannot share a slot”. Every pair is joined, so no two exams can be together, so you need three slots. That is obvious here with three exams. With three hundred it is not, and the useful part is that the same encoding answers both.

9 The pipeline



Two exits, and they are not symmetric — Section 5 again. The dashed boxes are plain text files. Nothing in this pipeline is a private binary format, which is why the Rust half can pick up where the Python half stopped.

10 Encoding it

One boolean per (exam, slot) pair: $x_{e,s}$ is true when exam e sits in slot s . Two kinds of constraint:

- Every exam gets exactly one slot: `exactly_one` over $x_{e,*}$.
- Clashing exams differ: for each clashing pair and each slot, $(\neg x_{a,s} \vee \neg x_{b,s})$.

"""One problem, end to end, showing every intermediate.

Three exams. Some students are enrolled in both Algebra and Biology, some in both Biology and Chemistry, some in both Algebra and Chemistry. An exam pair with a shared student cannot go in the same slot.

Two slots is impossible. Three works. We do both, and write out the files that

pass between the stages.

"""

```
import pathlib
from cdclkit import CNF, Encoder, solve, neg, write_dimacs, MemoryProof
import dratify

EXAMS = ["Algebra", "Biology", "Chemistry"]
CLASHES = [("Algebra", "Biology"), ("Biology", "Chemistry"), ("Algebra", "Chemistry")]
OUT = pathlib.Path(__file__).resolve().parent.parent / "build"
OUT.mkdir(exist_ok=True)

def build(n_slots):
    """Encode the timetable. Returns (formula, variable map)."""
    f = CNF()
    enc = Encoder(f)
    # x[e][s] is true when exam e sits in slot s
    x = {e: [enc.new_lit() for _ in range(n_slots)] for e in EXAMS}

    for e in EXAMS:
        enc.exactly_one(x[e]) # every exam gets one slot
    for a, b in CLASHES:
        for s in range(n_slots): # clashing exams differ
            enc.add([neg(x[a][s]), neg(x[b][s])])
    return f, x

# ----- two slots
f2, _ = build(2)
proof = MemoryProof()
sat2, _ = solve(f2, proof=proof)

with open(OUT / "timetable2.cnf", "w") as fh:
    write_dimacs(f2, fh)
(OUT / "timetable2.drat").write_text(proof.to_text())

print(f"2 slots: {'SAT' if sat2 else 'UNSAT'}")
    f"    ({f2.nvars} vars, {len(f2.clauses)} clauses)")
print(f"    proof: {proof.n_add} additions, {proof.n_del} deletions")
print(f"    checked: {dratify.check_proof(f2, proof).ok}")
print(f"    wrote build/timetable2.cnf and build/timetable2.drat")

# ----- three slots
f3, x3 = build(3)
sat3, model = solve(f3)
print()
print(f"3 slots: {'SAT' if sat3 else 'UNSAT'}")
    f"    ({f3.nvars} vars, {len(f3.clauses)} clauses)")
for e in EXAMS:
    slot = next(s for s in range(3) if model[x3[e][s]] >> 1)
```

```
print(f" {e:<10} -> slot {slot + 1}")
```

```
2 slots: UNSAT (6 vars, 12 clauses)
  proof: 2 additions, 0 deletions
  checked: True
  wrote build/timetable2.cnf and build/timetable2.drat

3 slots: SAT (9 vars, 21 clauses)
  Algebra    -> slot 3
  Biology    -> slot 2
  Chemistry  -> slot 1
```

11 What actually got produced

The formula

Six variables for three exams over two slots. Written out, the whole thing:

```
p cnf 6 12
1 2 0      <- Algebra in slot 1 or 2
-1 -2 0    <- ... but not both
3 4 0      <- Biology
-3 -4 0
5 6 0      <- Chemistry
-5 -6 0
-1 -3 0    <- Algebra and Biology not both in slot 1
-2 -4 0    <- ... nor both in slot 2
-3 -5 0    <- Biology and Chemistry
-4 -6 0
-1 -5 0    <- Algebra and Chemistry
-2 -6 0
```

Variable $2e + s$ is exam e in slot s . The annotations are mine; the file has none. This is what every SAT tool in the world consumes.

The proof

```
-2 0
0
```

Two lines. Read it as: “Algebra is not in slot 2” follows from the formula, and once you have that, the empty clause follows.

Check the first by hand with the RUP rule. Assume the opposite — Algebra is in slot 2, so variable 2 is true. Then $-1 -2 0$ forces variable 1 false, and $1 2 0$ is already satisfied. Follow the clashes: $-2 -4 0$ forces 4 false, so $3 4 0$ forces 3 true; $-2 -6 0$ forces 6 false, so $5 6 0$ forces 5 true. Now 3 and 5 are both true — Biology and Chemistry both in slot 1 — and $-3 -5 0$ forbids exactly that. Contradiction, so the clause holds.

That is the entire proof system. A checker does this mechanically, for every line, and reports whether the empty clause is ever reached.

12 Picking it up from Rust

The two files are the whole interface. Nothing else passes between the processes, and the checking side never sees the solver.

```

///! Read the DIMACS and DRAT files Python wrote, and verify the refutation.
///!
///! This is the realistic shape: some other program produced the proof, and
///! this process has to decide whether to believe it. The two files are the
///! entire interface between them.

use dratify::{Checker, Lit};
use std::fs;

/// DIMACS literal (1-based, signed) to the internal doubled index.
fn lit(d: i32) -> Lit {
    let v = (d.unsigned_abs() - 1) as Lit;
    (v << 1) | if d < 0 { 1 } else { 0 }
}

/// Parse DIMACS CNF. Returns (number of variables, clauses).
fn parse_cnf(text: &str) -> (u32, Vec<Vec<Lit>> > {
    let mut nvars = 0;
    let mut clauses = Vec::new();
    for line in text.lines() {
        let line = line.trim();
        if line.is_empty() || line.starts_with('c') {
            continue;
        }
        if let Some(rest) = line.strip_prefix("p cnf ") {
            nvars = rest.split_whitespace().next().unwrap().parse().unwrap();
            continue;
        }
        clauses.push(line.split_whitespace()
            .map(|t| t.parse::<i32>().unwrap())
            .take_while(|&d| d != 0)
            .map(lit)
            .collect());
    }
    (nvars, clauses)
}

/// Parse a text DRAT proof into (is_deletion, literals) steps.
fn parse_drat(text: &str) -> Vec<(bool, Vec<Lit>) > {
    text.lines()
        .map(str::trim)
        .filter(|l| !l.is_empty())
        .map(|line| {
            let (del, body) = match line.strip_prefix("d ") {
                Some(rest) => (true, rest),
                None => (false, line),
            };
            (del, body.split_whitespace()
                .map(|t| t.parse::<i32>().unwrap())
                .take_while(|&d| d != 0)
                .map(lit)
                .collect())
        })
}

```

```

        .collect()
    }

    fn main() {
        let dir = concat!(env!("CARGO_MANIFEST_DIR"), "/../../build");
        let cnf = fs::read_to_string(format!("{dir}/timetable2.cnf"))
            .expect("run code/ex6_pipeline.py first to generate the files");
        let drat = fs::read_to_string(format!("{dir}/timetable2.drat")).unwrap();

        let (nvars, clauses) = parse_cnf(&cnf);
        let proof = parse_drat(&drat);
        println!("formula: {} variables, {} clauses", nvars, clauses.len());
        println!("proof:    {} steps", proof.len());

        let r = Checker::new(nvars, &clauses, true, true).check(&proof);
        println!();
        println!("verified:    {}", r.ok);
        println!("empty clause: {}", r.reached_empty);
        println!("by RUP:      {}", r.rup_steps);
        println!();
        println!("=> two exam slots really are impossible, and this process");
        println!("    confirmed it without trusting the solver that said so.");
    }
}

```

```

formula: 6 variables, 12 clauses
proof:    2 steps

verified:    true
empty clause: true
by RUP:      2

=> two exam slots really are impossible, and this process
    confirmed it without trusting the solver that said so.

```

The parsing is a third of that file, and it is the honest cost of a text interchange format. What you get for it: the program deciding whether to believe the answer shares nothing with the program that produced it.

Part III

Python

13 Installing

```
pip install cdclkit          # solver, encodings, modelling layer
pip install "cdclkit[native]" # the same, plus a Rust engine ~18x faster
```

`cdclkit` depends only on `dratify`, the proof checker, which has no dependencies of its own. Nothing third-party is pulled in.

14 Your first formula

```
"""Smallest possible: write a formula by hand, ask whether it can be
   satisfied."""
from cdclkit import parse_dimacs, solve

# (a OR b) AND (NOT a OR b) AND (a OR NOT b)
# DIMACS: variables are 1, 2, ...; a negative number is a negated variable;
# each clause ends with 0.
text = """
p cnf 2 3
 1  2 0
-1  2 0
 1 -2 0
"""

formula = parse_dimacs(text)
print("variables:", formula.nvars, " clauses:", len(formula.clauses))

sat, model = solve(formula)
print("satisfiable:", sat)
if sat:
    # model is indexed by variable number - 1
    print("a =", model[0], " b =", model[1])
```

```
variables: 2  clauses: 3
satisfiable: True
a = True  b = True
```

`solve` returns a *tuple*. Note the shape — if `solve(f)`: is always true, because a two-element tuple is always truthy. Unpack it.

15 Unsatisfiable, and proving it

Forbid all four combinations of two variables and nothing is left.

```

"""An unsatisfiable formula, and the proof that it is."""
from cdclkit import parse_dimacs, solve, MemoryProof
import dratify

# All four combinations of two variables are forbidden, so nothing is left.
formula = parse_dimacs("""
p cnf 2 4
 1 2 0
 1 -2 0
-1 2 0
-1 -2 0
""")

proof = MemoryProof()
sat, model = solve(formula, proof=proof)
print("satisfiable:", sat)
print("proof steps:", proof.n_add, "additions,", proof.n_del, "deletions")

result = dratify.check_proof(formula, proof)
print("proof verified by an independent checker:", result.ok)
print("empty clause derived:", result.reached_empty)

```

```

satisfiable: False
proof steps: 2 additions, 0 deletions
proof verified by an independent checker: True
empty clause derived: True

```

This is Section 5 made concrete. The solver said “no” and handed over a proof; `dratify` — a different package, sharing no solver code — replayed it and agreed. You did not have to trust the solver.

16 Constraints, without writing clauses

Real problems are not phrased as clauses. They are phrased as “exactly one of these”, “at most three of those”. The Encoder turns those into clauses for you.

```

"""Constraints people actually have, without writing clauses by hand.

Colour a triangle with 2 colours. Each node needs exactly one colour, and
adjacent nodes must differ. A triangle needs 3 colours, so this is UNSAT.
"""
from cdclkit import CNF, Encoder, solve, neg

for n_colours in (2, 3):
    f = CNF()
    enc = Encoder(f)

    nodes = ["x", "y", "z"]
    edges = [("x", "y"), ("y", "z"), ("x", "z")]

    # colour[node][c] is true when `node` has colour c
    colour = {n: [enc.new_lit() for _ in range(n_colours)] for n in nodes}

    for n in nodes:
        enc.exactly_one(colour[n])          # one colour per node
    for a, b in edges:

```

```

    for c in range(n_colours):
        # never the same colour on an
        edge
        # note neg(), not -x: internal literals are non-negative
        enc.add([neg(colour[a][c]), neg(colour[b][c])])

    sat, model = solve(f)
    print(f"{n_colours} colours: {'SAT' if sat else 'UNSAT'}"
          f"    ({f.nvars} vars, {len(f.clauses)} clauses)")

```

```

2 colours: UNSAT    (6 vars, 12 clauses)
3 colours: SAT     (9 vars, 21 clauses)

```

A triangle needs three colours, and the solver proves two are impossible.

One gotcha. Internal literals are non-negative — a variable v is $2v$ positive and $2v+1$ negated. So negation is `neg(x)`, not `-x`. This differs from DIMACS on purpose: it makes array indexing direct in the solver's hot loops. Use `from_dimacs` when converting signed input.

The encoder also offers `at_most_one`, `at_least_k`, `exactly_k`, pseudo-boolean constraints (`assert_pb_leq`), and Tseitin gates (`and_gate`, `or_gate`, `xor_gate`, `ite`).

17 The modelling layer

For most problems you should not touch literals at all. The modelling layer gives you finite-domain integers.

```

"""The modelling layer: integers and all-different, no clauses in sight.

A 4x4 Sudoku. Each cell holds 1..4; rows, columns and 2x2 boxes are
all-different.
"""
from cdclkit.model import Model

PUZZLE = [
    [1, 0, 0, 0],
    [0, 0, 3, 0],
    [0, 4, 0, 0],
    [0, 0, 0, 2],
]

m = Model()
cell = [[m.int_var(range(1, 5), f"c{r}{c}") for c in range(4)] for r in range(4)]

for r in range(4):
    m.all_different(cell[r])
    # rows
for c in range(4):
    m.all_different([cell[r][c] for r in range(4)])
    # columns
for br in (0, 2):
    for bc in (0, 2):
        m.all_different([cell[br + i][bc + j] for i in range(2) for j in range(2)])
        # 2x2 boxes

for r in range(4):
    for c in range(4):
        if PUZZLE[r][c]:

```

```

        m.add(cell[r][c] == PUZZLE[r][c])                # the givens

sol = m.solve()
print("solved:", sol is not None)
for r in range(4):
    print("    ", " ".join(str(sol.value(cell[r][c])) for c in range(4)))

```

```

solved: True
  1 3 2 4
  4 2 3 1
  2 4 1 3
  3 1 4 2

```

`int_var` takes a domain, `all_different` does what it says, and `cell == 3` builds a constraint. Underneath it is still CNF and still the same solver — but you never see a clause.

This is the layer to reach for first. Drop to the encoder when you need a constraint it does not offer, and to raw clauses almost never.

18 Why is it unsatisfiable?

When a model has no solution, the useful question is which constraints conflict. A **minimal unsatisfiable subset** answers it.

```

"""Why is it unsatisfiable? Ask for the part that matters.

A minimal unsatisfiable subset (MUS) is a subset of the clauses that is still
unsatisfiable, and stops being so if you drop any one of them.
"""
from cdclkit import parse_dimacs, mus

# Clause 4 is irrelevant to the contradiction between 1, 2 and 3.
formula = parse_dimacs("""
p cnf 3 4
  1 0
-1 2 0
-2 0
  3 0
""")

core = mus(formula)
print("clause indices in the minimal unsatisfiable subset:", core)
print("dropping any one of them makes it satisfiable again")

```

```

clause indices in the minimal unsatisfiable subset: [0, 1, 2]
dropping any one of them makes it satisfiable again

```

Clause 3 is irrelevant to the contradiction and the MUS excludes it. On a scheduling problem with a thousand constraints, this is the difference between “infeasible” and “these four requirements cannot all hold”.

19 The command line

```

cdclkit solve problem.cnf --self-check --check-model

```

Solves, and then: if SAT, re-evaluates the model against the original formula clause by clause; if UNSAT, emits a DRAT proof and has it independently replayed. Exit codes follow the SAT competition convention — 10 for satisfiable, 20 for unsatisfiable — which is unusual enough to surprise a shell script expecting 0.

20 What the Rust crate is for

If you want to *solve* from Rust, use one of the established solvers. If you want to *check*, this is a dependency-free crate that does RUP and RAT properly.

The crate is `dratify`; the source for both it and the Python package is `dratify`.

21 Literals

DIMACS	meaning	internal
1	a	$\emptyset$
-1	$\neg a$	1
2	b	2
-2	$\neg b$	3

Variable v becomes $2v$ positive, $2v+1$ negated. Negation is a single XOR with 1, which is why the encoding exists.

22 Checking a proof

17

```

        vec![lit(1), lit(2)],
        vec![lit(1), lit(-2)],
        vec![lit(-1), lit(2)],
        vec![lit(-1), lit(-2)],
    ];

    // A DRAT proof: derive "a", then the empty clause. Each step is
    // (is_deletion, literals).
    let proof: Vec<(bool, Vec<Lit>)> = vec![
        (false, vec![lit(1)]),
        (false, vec![]),
    ];

    let mut checker = Checker::new(2, &formula, true, true);
    let r = checker.check(&proof);

    println!("verified:      {}", r.ok);
    println!("empty clause:  {}", r.reached_empty);
    println!("steps checked: {} ({} by RUP, {} by RAT)", r.steps, r.rup_steps,
        , r.rat_steps);

    // Now a dishonest proof: claim the empty clause with nothing to back it.
    let bogus: Vec<(bool, Vec<Lit>)> = vec![(false, vec![])];
    let mut c2 = Checker::new(2, &vec![vec![lit(1)]], true, true);
    let bad = c2.check(&bogus);
    println!();
    println!("bogus refutation accepted: {}", bad.ok);
    println!("reason: {}", bad.reason);
}

```

```

verified:      true
empty clause:  true
steps checked: 2 (2 by RUP, 0 by RAT)

bogus refutation accepted: false
reason: clause is neither RUP nor RAT with respect to the clauses available
        at this point

```

The second half matters more than the first. A checker that accepts everything would pass every test that only feeds it valid proofs. Any checker you rely on should be tested on proofs that are *wrong*, and should reject them for the right reason.

23 What the result tells you

check returns a CheckResult:

ok	the proof is a valid refutation
reached_empty	the empty clause was actually derived
steps, rup_steps, rat_steps	what was checked, and how
failed_step, failed_clause	where it went wrong, if it did
reason	why, in words

Note that ok and reached_empty are separate. A proof whose every step verifies but which never derives the empty clause is not a refutation — it is a valid but incomplete derivation. Conflating the two is how a truncated proof gets accepted.

24 Two implementations on purpose

The Python package ships its own pure-Python checker, and this crate is an independent implementation of the same rules. Neither is the “real” one.

For a proof checker that is the entire point: agreement between independent implementations is the evidence. They are differentially tested against each other on rejections as well as acceptances, which is the half that catches a checker that has quietly started saying yes to everything.

From Python you can have both. `pip install "cdclkit[native]"` installs wheels that register the Rust checker with `dratify`, and `check_proof(..., engine="python")` versus `engine="native"` then runs each.

Where to go next

- `examples/` in the `cdclkit` repository — Sudoku with a uniqueness proof, the zebra puzzle, graph colouring, bounded model checking, circuit equivalence.
- `docs/ALGORITHMS.md` — the mathematics, from resolution through first-UIP learning, LBD, DRAT and preprocessing, including a section on what is deliberately absent.
- `BENCHMARKS.md` — how fast this is, and the caveats that belong with every number in it.
- For industrial-scale problems, `python-sat` bundles `kissat` and `CaDiCaL` as binary wheels, and is much faster. What it does not do is check the proofs it can emit — see Section 2.
- `drat-trim` is the reference DRAT checker, in C. `dratify` agrees with it on every proof they have both been given.